

Extrinsic Prior for Simple and Efficient Bayesian Modeling with Parameter Constraints

Abstract: Parameter constraints are very common in statistical models. Examples include linear inequality, parameter ordering, monotonicity, orthogonality, etc. Bayesian approach is useful for uncertainty quantification in the constrained space. Although specific solutions have been made for different constraints, it is challenging to incorporate them in advanced applications, such as modeling with non-parametric assumption or high-dimensional data. In this paper, we propose a simple and general solution by first replacing constraints with strongly informative prior. Through this *extrinsic* prior, the parameters are relaxed to a less restrictive space, where conventional tools such as Hamiltonian Monte Carlo are utilized to obtain approximate posterior. Then these posteriors can be easily projected back to the constrained space for exact solution. This approach is very efficient and applicable to a wide range of problems with equality and inequality constraints. The generality allows more families of prior to be chosen for the constrained parameters, and simplifies the adoption of multiple constraints for desired property such as identifiability. Theory is developed and novel statistical applications under constraints are illustrated.

KEY WORDS: Constraint relaxation; Space embedding; Monotone Dirichlet; Orthogonal Gaussian processes; Posterior mixing; Projected Markov chain

1 Introduction

Constraints are very common in modern statistical modeling. For example, functional data analysis often impose certain shape constraint such as monotonicity or convexity on curves (Kelly and Rice, 1990); matrix and tensor decomposition utilize orthonormality for better identifiability (Uschmajew, 2010); many manifolds such as simplex can be considered as sub-manifolds of a Euclidean space embedded via certain constraints.

Constraints can cause substantial modeling difficulty. When data are in constrained space, parameters can enter the likelihood via an integral without closed-form, commonly known as “doubly intractable” problem. Successful solutions have been proposed to address this issue (Murray et al., 2012; Rao et al., 2016). When parameters are in constrained space, challenges often arise in the difficulty for estimation under constraint. Frequentist optimization literature often relies on Lagrange and Karush-Kuhn-Tucker multipliers for point

estimate under equality and inequality constraints (Boyd and Vandenberghe, 2004). However, the uncertainty quantification is difficult since conventional asymptotic result on variance estimation often no longer hold in constrained space. Bayesian approach is more appropriate for this purpose.

There have been a wide range of solutions developed for specific constraints. One strategy involves using constrained prior with posterior that can be conveniently sampled. For example, to model orthornormal matrices on the Stiefel manifold, Bingham-von-Mises-Fisher distribution (Khatri and Mardia, 1977; Hoff, 2009) is a parametric family with a closed-form posterior in matrix and tensor decomposition. Lin et al. (2016) extends the flexibility of matrix von-Mises-Fisher distribution via non-parametric approach. Another strategy is to bypass the constraint via re-parameterization. The famous example is the stick-breaking construction for Dirichlet distribution and process. The re-parameterization essentially utilizes the coordinate system of the simplex, and circumvents the norm constraint on the probability vertices. As these methods directly satisfy the constraint requirement, we refer them as the intrinsic approaches. Despite the success of intrinsic approaches, the posterior can quickly become very involved to sample under slightly more advanced model or complicated data. For example, in modeling population of undirected networks, the symmetry in each network disrupts the closed form of posterior in orthogonal tensor decomposition (Hoff et al., 2016), demanding new rejection sampling algorithm to be developed. As another example, additional structure (such as ordering) on the probability simplex would disrupt the simple form of stick-breaking posterior.

These drawbacks have motivated the development of extrinsic approaches. The key idea is to first sample the proposal freely in a conventional space (such as Euclidean space), then transform it back to the constrained space. One early work can be traced back to Gelfand et al. (1992), who used Gibbs sampling to first generate proposal in unrestricted region, then only accepting those inside the constraint space. One critical issue is that unrestricted proposal can have significant mass outside the constraint region, resulting in a high rejection rate. Replacing rejection sampling, Lin and Dunson (2014) and Lin et al. (2016) utilize projection to map the unconstrained posterior into the constrained space and obtain monotonicity and manifold-valued regression. These specialized cases seem to work well, but the suitable projection are often elusive and there is a clear lack of general and simple approach.

In this paper, we propose a general extrinsic approach, by parameterizing constraints as a limiting case of strongly informative prior. We refer them as extrinsic priors. We then relax the effective support of the prior to a neighborhood of constraint space, obtaining posterior via efficient tools such as conventional Hamiltonian Monte Carlo (HMC). These posteriors are approximate to the canonical formulation, with approximation error bounded during the prior specification. The imperfection of approximation can be corrected with a simple projection and a Metropolis-Hastings step with high acceptance probability, leading to a Markov chain corresponding to the exact formulation. Compared to other manifold based methods such as Riemannian

and geodesic HMC (Girolami and Calderhead, 2011; Byrne and Girolami, 2013), our approach is efficient in computation and simple to implement via highly automatic software like STAN. The simplicity enables a larger spectrum of prior to be chosen and more free adaption of constraints in modeling. Theoretic studies are conducted and original models are shown in simulations and data application.

2 Method

We consider a parameter θ in a constrained space \mathcal{D} . The space \mathcal{D} can be high- or infinite-dimensional. Letting this space be equipped with a σ -field \mathcal{B} , the standard Bayesian approach assigns a prior for θ in \mathcal{D} , based on a density $\pi_{0,\mathcal{D}}(\theta)$ in a separable space $(\mathcal{D}, \mathcal{B})$. In intrinsic approaches, priors are chosen for computational conveniences so that the posterior can be easily sampled strictly inside \mathcal{D} . Clearly, the choices of priors and constraints one can impose are very limited. Instead, we consider extrinsic approaches by estimating θ in the larger space \mathcal{R} where $\mathcal{D} \in \mathcal{R}$. We first provide a probabilistic justification.

Assuming $\pi_{0,\mathcal{D}}(\theta)$ is proper $\int_{\mathcal{D}} \pi_{0,\mathcal{D}}(\theta) d\theta < \infty$, then this prior can be viewed as a conditional density, based on another density $\pi_{0,\mathcal{R}}(\theta)$ in $(\mathcal{R}, \mathcal{C})$ with \mathcal{C} as the σ -field of \mathcal{R} :

$$\pi_{0,\mathcal{D}}(\theta) = \pi_{0,\mathcal{R}}(\theta \mid \theta \in \mathcal{D}) = \frac{\pi_{0,\mathcal{R}}(\theta) \mathbb{1}_{\theta \in \mathcal{D}}}{\int_{\mathcal{D}} \pi_{0,\mathcal{R}}(\theta) d\theta}. \quad (1)$$

where $\mathbb{1}_{\theta \in \mathcal{D}} = 1$ when $\theta \in \mathcal{D}$, 0 otherwise. Note as long as $\pi_{0,\mathcal{R}}(\theta) \mathbb{1}_{\theta \in \mathcal{D}}$ is proper, $\pi_{0,\mathcal{R}}(\theta)$ can be improper. Letting $L(\theta; y)$ be the likelihood function and y be the observed data, the posterior can be obtained via:

$$\pi(\theta \mid y, \theta \in \mathcal{D}) = \frac{L(\theta; y) \pi_{0,\mathcal{D}}(\theta)}{\int_{\mathcal{D}} L(\theta; y) \pi_{0,\mathcal{D}}(\theta) d\theta} = \frac{L(\theta; y) \pi_{0,\mathcal{R}}(\theta) \mathbb{1}_{\theta \in \mathcal{D}}}{\int_{\mathcal{D}} L(\theta; y) \pi_{0,\mathcal{R}}(\theta) d\theta}, \quad (2)$$

where the last equality holds because $\int_{\mathcal{D}} \pi_{0,\mathcal{R}}(\theta) d\theta$ is a finite constant.

2.1 Extrinsic Prior for Constraints

One obvious extrinsic approach utilizing (2) is to first generate proposal in \mathcal{R} based on $L(\theta; y) \pi_{0,\mathcal{R}}(\theta)$ (assuming it is proper), then accepting it when it falls in \mathcal{D} (Gelfand et al., 1992). However, when the probability $\pi(\theta \in \mathcal{D} \mid y) / \pi(\theta \in \mathcal{R} \setminus \mathcal{D} \mid y) \approx 0$, this leads to significant amount of rejections. This is very common in equality constraint.

We propose a different strategy. Instead of ignoring $\mathbb{1}_{\theta \in \mathcal{D}}$ in the first step, we approximate it with an additional strongly informative prior $\mathcal{E}(\theta)$. The prior has its support in \mathcal{R} , but concentrated around \mathcal{D} . When $\theta \in \mathcal{D}$, $\mathcal{E}(\theta)$ is constant; $\theta \notin \mathcal{D}$, $\mathcal{E}(\theta)$ quickly drops to 0. Then one can first obtain approximate posterior based on density proportional to $L(\theta; y) \pi_{0,\mathcal{R}}(\theta) \mathcal{E}(\theta)$ (the conditions for posterior propriety is postponed to the theory section).

In this paper, we focus on the embedding of \mathcal{D} in \mathcal{R} via equality and inequality constraints, although other types of constraints can be incorporated similarly. There are m equalities and l inequalities, leading to $\mathcal{D} = \{\theta \in \mathcal{R} : E_k(\theta) = 0 \text{ for } k = 1, \dots, m, \quad G_{k'}(\theta) \leq 0 \text{ for } k' = 1, \dots, l\}$, where $E_k(\cdot)$ and $G_{k'}(\cdot)$ are functions that map from \mathcal{R} to real line \mathbb{R} . Then the indicator function is $\mathbb{1}_{\theta \in \mathcal{D}} = \prod_k \mathbb{1}_{E_k(\theta)=0} \cdot \prod_{k'} \mathbb{1}_{G_{k'}(\theta) \leq 0}$.

We now replace the indicator functions with $\mathcal{E}(\theta)$, represented as a product of $(m + l)$ kernel functions $K(\cdot)$, leading to approximate posterior:

$$\begin{aligned} \pi_K(\theta | y) &\propto L(\theta; y) \pi_{0, \mathcal{R}}(\theta) \mathcal{E}(\theta) \\ &\propto L(\theta; y) \pi_{0, \mathcal{R}}(\theta) \cdot \prod_{k=1}^m K_{1,k}(|E_k(\theta)|) \cdot \prod_{k'=1}^l K_{2,k'}((G_{k'}(\theta))_+) \end{aligned} \quad (3)$$

where $(x)_+ = x$ if $x > 0$, 0 if $x \leq 0$. The posterior $\pi_K(\theta | y)$ is an approximation to $\pi(\theta | y)$ in (2). We will now refer $\pi_K(\theta | y)$ as “extrinsic posterior”. The functions $|E_k(\theta)| \in [0, \infty)$ or $(G_{k'}(\theta))_+ \in [0, \infty)$ represent the amount of relaxation of each constraint, where 0 represents no relaxation. Each kernel $K_{i,k}$ satisfies $K_{i,k}(0) = 1$; the tolerable amount of relaxation is controlled by a hyper-parameter $\lambda_{i,k}$. When $\lambda_{i,k} \rightarrow \infty$, the kernel becomes a point mass at 0. Therefore, (2) is a limiting case of (3). For example, one simple and useful kernel is the truncated Gaussian $K_{i,k}(x) = \exp(-\lambda_{i,k}x^2) \mathbb{1}_{x < \varepsilon}$.

Instead of taking infinite values for λ 's, we assign large but finite ones. This gives rise to a continuous relaxation of the sharp boundary of the indicator function. The relaxation allows the posterior θ to be easily sampled in \mathcal{R} under the guidance of the constraints. For example, one can carry out conventional HMC for constrained parameters directly in Euclidean space. At the same time, since posteriors are generated in a tight neighborhood of \mathcal{D} , they can be easily projected back to \mathcal{D} as proposals in a correcting Metropolis-Hastings step.

2.2 Control of Constraint Relaxation

In extrinsic posterior (3), when $\theta \in \mathcal{D}$, the density is the same as (2), up to a constant difference. However, since we induce positive support in $\mathcal{R} \setminus \mathcal{D}$, it is worth studying how the approximate posterior are distributed relatively to the space \mathcal{D} . This is reflected in the posterior distribution of the constraint relaxation $|E_k(\theta)|$ and $(G_{k'}(\theta))_+$.

The constraint relaxation can be controlled via a tight prior support near 0 for each kernel. That is $\int_{x < \varepsilon} \mathcal{C}_{i,k}(x) dx = 1$, with $\mathcal{C}_{i,k}(x) = K_{i,k}(x) / \int_{\mathcal{R}} K_{i,k}(x) dx$. The pre-specified constant ε represents the element-wise tolerance for violating each constraint. The bounded prior support allows us to theoretically control the posterior approximation error. As $\mathcal{E}(\theta) \propto \prod_{i,k} \mathcal{C}_{i,k}(x)$ is the joint extrinsic prior density, since $\pi_K(\theta | y) \ll \mathcal{E}(\theta)$, the posterior for each constraint relaxation is bounded in $[0, \varepsilon)$ with probability 1.

In practice, one may wish to utilize a kernel $K_{i,k}^*(x)$ with unbounded support on $[0, \infty)$ for computing conveniences. To adopt them, one can first choose $\lambda_{i,k}$ to have $\int_{x < \varepsilon} K_{i,k}^*(x) / (\int_{\mathcal{R}} K^*(x) dx) = 1 - \eta$ with η small, then apply truncation $K_{i,k}(x) = K_{i,k}^*(x) \mathbb{1}_{x < \varepsilon}$ to ensure $x < \varepsilon$ almost surely. In most cases, the truncation is just nominal for a theoretic guarantee; in computation it can be satisfied automatically with high probability. For example, in Gaussian kernel $\exp(-\lambda x^2)$ setting $\lambda = \frac{1}{2(\varepsilon/4)^2}$ ensures $x < \varepsilon$ with probability 0.99993 apriori; for posterior sampling, one can first do an untruncated sampling, then reject those $x < \varepsilon$, which is likely rare due to the small prior probability.

To illustrate the approximation of extrinsic prior and control of constraint relaxation, we consider a simple example of generating a truncated Gaussian distribution $\theta \mid y \sim \text{No}_{(\alpha, \beta)}(0, 1)$, with mean 0 and variance 1 and truncation $\theta \in (\alpha, \beta)$. The exact and extrinsic posterior densities are:

$$\pi(\theta \mid y) \propto \exp\left(-\frac{\theta^2}{2}\right) \mathbb{1}_{\theta \in (\alpha, \beta)}, \quad \pi_K(\theta \mid y) \propto \exp\left(-\frac{\theta^2}{2}\right) K((\alpha - \theta)_+) K((\theta - \beta)_+).$$

with $K(x) = \exp(-\lambda x^2) \mathbb{1}_{x < 4/\sqrt{2\lambda}}$. We set $(\alpha, \beta) = (1, 2)$. Figure 1 plots the unnormalized densities under the exact posterior and approximation with different λ 's. The approximate densities inside $\mathcal{D} = (1, 2)$ are the same as the exact one, up to a constant difference due to normalization. Outside \mathcal{D} , the larger λ is associated with more rapid decline of density and therefore smaller tolerance for constraint relaxation.




Figure 1: Unnormalized densities for truncated normal $\text{No}_{(1,2)}(0, 1)$, under exact and approximating densities. The exact density abruptly drops to 0 on the two boundaries, while the approximating ones drop continuously. In the approximation, larger λ is associated with lower tolerance for constraint relaxation $((1 - \theta)_+$ and $(\theta - 2)_+$). All densities inside $(1, 2)$ are the same up to a constant difference.

Although it is tempting to always induce almost 0 relaxation with very large λ , in heavily constrained models such as the ones with equality constraint, the narrow distribution width in \mathcal{R} will cause a detrimental effect in algorithms such as Hamiltonian Monte Carlo. In those cases, it is rather useful to have a slightly larger relaxation, then use projection to correct the imperfection. We will illustrate this in the next two sections.

2.3 Posterior Sampling for Extrinsic Posterior

Extrinsic posterior is the approximation to the ones under exact formulation. As it is defined on a less restrictive space \mathcal{R} , it can be sampled easily. The traditional sampling tools such as slice sampling, adaptive Metropolis-Hastings can be utilized. In this section, we present the sampling algorithm using Hamiltonian Monte Carlo (HMC), due to its high-level automation aided by software and excellent performance in convergence and posterior mixing. Various adaptive algorithms such as Hoffman and Gelman (2014) have been developed for making the new state less correlated with the current state.

In using the conventional HMC, we assume θ is d -dimensional and the space \mathcal{R} is an Euclidean space \mathbb{R}^d and the constraint functions $E_k(\theta)$'s and $G_k(\theta)$'s are differentiable with respect to θ . We focus on the case where θ is continuous, although discrete extension is possible (Zhang et al., 2012).

HMC is essentially a data augmentation based MCMC. Using a latent variable named “veolicty” $p \in \mathbb{R}^d$, the negative log-posterior function based on (3) is

$$\begin{aligned} H(\theta, p) &= U(\theta) + M(p), \\ \text{where } U(\theta) &= -\log \{L(\theta; y)\pi_{0,\mathcal{R}}(\theta)\mathcal{E}(\theta)\}, \\ M(p) &= \frac{p'\Sigma^{-1}p}{2}, \end{aligned} \tag{4}$$

with Σ^{-1} a pre-specified positive definite matrix. Instead of using random walk or Gibbs sampling, HMC update θ and p via Hamiltonian dynamics, satisfying differential equations:

$$\begin{aligned} \frac{\partial \theta(t)}{\partial t} &= \frac{\partial H(\theta, p)}{\partial p} = \Sigma^{-1}p, \\ \frac{\partial p(t)}{\partial t} &= -\frac{\partial H(\theta, p)}{\partial \theta} = -\frac{\partial U(\theta)}{\partial \theta}. \end{aligned} \tag{5}$$

At the start of each iteration, the current state of θ is viewed as $\theta(0)$ with $p(0)$ randomly generated from $\text{No}(0, \Sigma)$. The solution to (5) yields $\theta(t)$ and $p(t)$. Since Hamiltonian system is symplectic, $H(\theta(t), p(t)) = H(\theta(0), p(0))$. However, in most cases, (5) lacks closed-form solution, one has to use discrete approximation, such as leap-frog algorithm (Neal, 2011):

$$\begin{aligned} p(T + \varepsilon/2) &= p(T) - \varepsilon/2 \frac{\partial U}{\partial \theta}(\theta(T)), \\ \theta(T + \varepsilon) &= \theta(T) + \varepsilon \Sigma^{-1}p(T + \varepsilon/2), \\ p(T + \varepsilon) &= p(T + \varepsilon/2) - \varepsilon/2 \frac{\partial U}{\partial \theta}(\theta(T + \varepsilon)), \end{aligned} \tag{6}$$

for $T = 0, \varepsilon, 2\varepsilon, \dots, (L-1)\varepsilon$, with ε known as the time step, and L as the total leap-frog steps within one iteration. Since this approximating update is reversible, an Metropolis-Hastings step is taken at the end to accept $\theta(t)$ with probability $1 \wedge \exp(-H(\theta(t), -p(t)) + H(\theta(0), p(0)))$ with $t = L\varepsilon$. Since constraints $\theta \in \mathcal{D}$ is now replaced by prior $\mathcal{E}(\theta)$, the derivative $\frac{\partial U}{\partial \theta}$ can be computed easily, allowing simple HMC to run in space \mathcal{R} .

We now focus on finding the optimal time step ε . The time step ε controls the stability of trajectory. When ε is too large, H grows exponentially with L , leading to very small acceptance rate. When ε is too small, it leads to wasteful computation in making very local update. Therefore, it is useful to set large ε near a stability bound. For simple system such as $U(\theta) = \theta^2/2\sigma^2$, one can write (6) as a linear transformation of $[\theta(T+\varepsilon), p(T+\varepsilon)]' = Q[\theta(T), p(T)]'$ (Q is a $2d \times 2d$ transition matrix), bounding the eigenvalues in Q below magnitude 1 determines the bound for ε . Since most systems involve nonlinear transition, analytical bound is not available, but one can empirically optimize ε to be close to this bound, by tuning for acceptance rate in the Metropolis-Hastings step. Specifically, given fixed L , one tunes ε so that the acceptance rate is close to but slightly below 1.

For multiple-dimensional θ with $\Sigma = I$, the stability bound is roughly determined by the width of distribution in the most constrained direction (Neal, 2011). To provide an intuition, we focus on $L = 1$. A stable trajectory would approximately preserve $U(\theta(\varepsilon)) + M(p(\varepsilon)) = U(\theta(0)) + M(p(0))$, since $M(p) = p'p/2 \geq 0$, $U(\theta(\varepsilon)) \leq U(\theta(0)) + M(p(0))$. With initial velocity $p(0) \sim N(0, I)$ and finite $U(\theta(0))$, a stable approximating trajectory should not move to position with infinite $U(\theta(t))$, which corresponds to 0 posterior density. However, in the leap-frog algorithm, since $p(0)$ is randomly generated in all direction of \mathbb{R}^d , ε needs to be small enough so that the discrete update $\theta(\varepsilon) = \theta(0) + \varepsilon p(0) - \varepsilon^2/2 \frac{\partial U}{\partial \theta}(\theta(0)) = \theta(0) + \varepsilon p(0) + O(\varepsilon^2)$ does not move $\theta(\varepsilon)$ outside the posterior support.

For constrained space \mathcal{D} formed by strong constraints such as equality, using a very large λ in $\mathcal{E}(\theta)$ would create a support with very narrow width along certain direction. This would force ε to be very small, leading to inefficient update. Therefore, for those constraints, it is rather useful to use a slightly smaller λ to induce more relaxation and allow more efficient Hamiltonian dynamics.

To illustrate, consider generating a random variable $\theta = (x_1, x_2)$ on a unit circle using von Mises-Fisher distribution, $\pi(\theta | y) \propto \exp(F'\theta)$ with $\theta'\theta = 1$. This is a simple example of a random variable constraint on a (2,1)-Stiefel manifold $\mathcal{D} = \mathcal{V}(2,1)$. We set $F = (1,1)$ to induce a distribution widely spreaded over the manifold, generating great amount of uncertainty for assessing the sampling efficiency. We use extrinsic prior proportional to $K(\theta) = \exp(-\lambda(\theta'\theta - 1)^2) \mathbb{1}_{|\theta'\theta - 1| < 0.1}$. Geometrically, this prior expands the posterior support from a circle to a ring, with its width determined by λ .

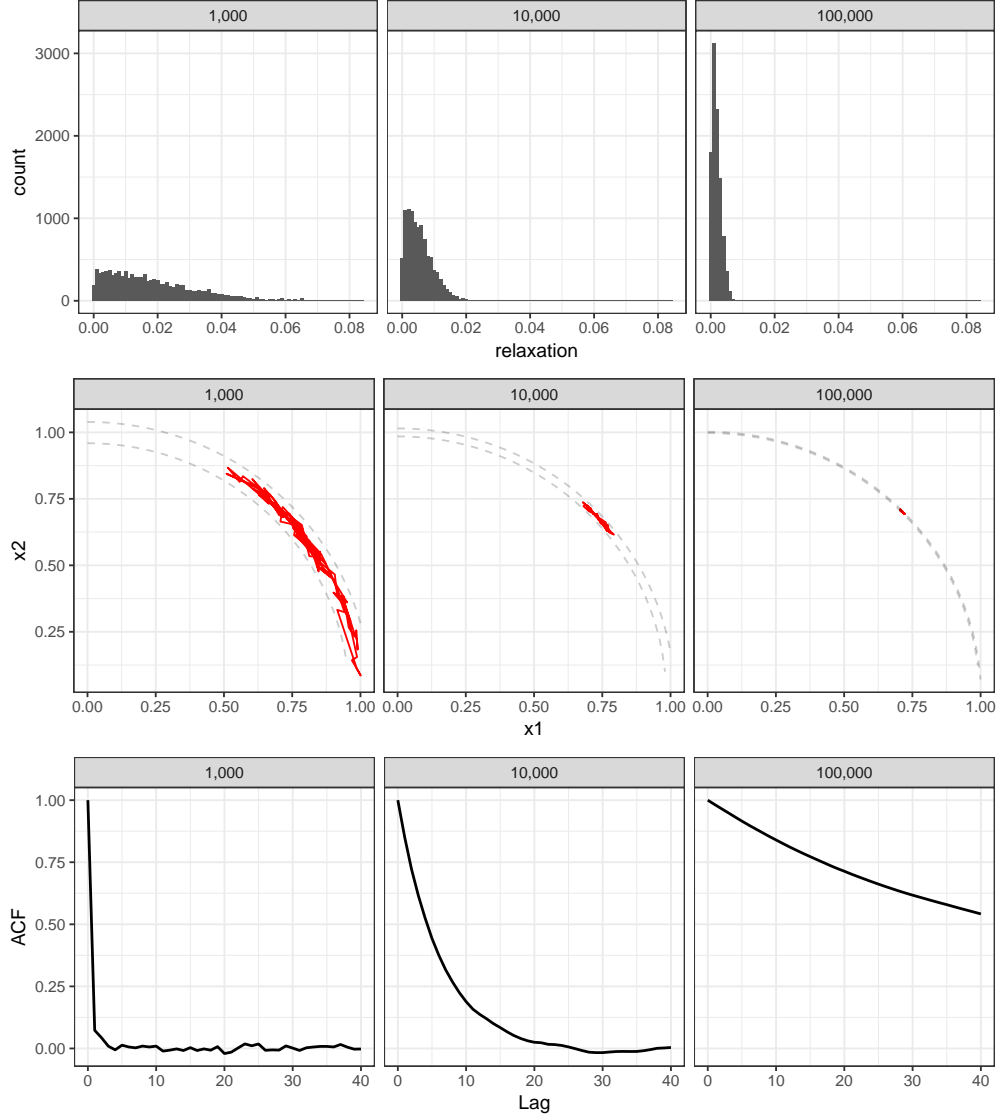


Figure 2: Sampling posterior from a von Mises–Fisher distribution on a unit circle, using HMC with extrinsic prior under $\lambda = 10^3, 10^4, 10^5$. Row 1 shows the posterior distribution of the constraint relaxation $|\theta'\theta - 1|$; Row 2 shows the path of 100 leap-frog steps; Row 3 shows the autocorrelation plot (ACF). Large λ gives very small constraint relaxation, but suffers from slow mixing due to inefficient local update; smaller λ increases the relaxation but results in excellent mixing.

We tested three different values of $\lambda = 10^3, 10^4, 10^5$. For each λ , we ran HMC for 10,000 iterations, with $L = 100$ leap-frog steps in each iteration. We set $\Sigma = \text{diag}(1, 1)$ in generating velocity p . During the initial 2,000 iterations, the leap-frog step size ε is tuned for an acceptance rate close to 0.8, then it is fixed during the remaining part of Markov chain. The last 5,000 iterations are used as posterior samples. Figure 2 plots the posterior distribution of constraint relaxation $|\theta'\theta - 1|$, the sampling path and the autocorrelation function (ACF) for each Markov chain. Very large $\lambda = 10^5$ has much less constraint relaxation; however, due to the small ring width, the Hamiltonian dynamics has to use small ε and can only explore local space for each 100 steps. This results in a very slow mixing (large autocorrelation even at 40 lags). On the other

hand, smaller $\lambda = 10^3$ has slightly larger constraint relaxation, but allows much more efficient exploration of the space and excellent mixing performance. We find that $\lambda = 10^3$ is a good empirical value for all the equality constraints in this paper.

2.4 Correcting Projection to Constraint Space

The Markov chain produced by HMC is geometrically ergodic under very general conditions (Livingstone et al., 2016). With the extrinsic posterior $\pi_K(\theta | y)$ as approximation to (2), one may be interested in further obtaining exact posterior in \mathcal{D} , likely for two reasons: (i) to strictly uphold the constraints; (ii) to ease the strict relaxation control on extrinsic prior.

Letting θ^* be a random sample collected based on $\pi_K(\theta | y)$, there exists deterministic projection $P : \mathcal{R} \rightarrow \mathcal{D}$ and obtain $\theta_{\mathcal{D}}^* = P(\theta^*)$. Using this as proposal machinery, one can construct another Markov chain with $\pi(\theta_{\mathcal{D}} | y)$ as the target distribution. Letting the current state be $\theta_{\mathcal{D}} = P(\theta)$, we generate proposal $\theta_{\mathcal{D}}^* = P(\theta^*)$ and accept it with probability:

$$1 \wedge \frac{\pi(\theta_{\mathcal{D}}^* | y) \pi_K(\theta | y)}{\pi(\theta_{\mathcal{D}} | y) \pi_K(\theta^* | y)} = 1 \wedge \frac{L(\theta_{\mathcal{D}}^*; y) \pi_{0, \mathcal{R}}(\theta_{\mathcal{D}}^*) \cdot L(\theta; y) \pi_{0, \mathcal{R}}(\theta) \mathcal{E}(\theta)}{L(\theta_{\mathcal{D}}; y) \pi_{0, \mathcal{R}}(\theta_{\mathcal{D}}) \cdot L(\theta^*; y) \pi_{0, \mathcal{R}}(\theta^*) \mathcal{E}(\theta^*)}. \quad (7)$$

The remaining task is then to optimize the projection with respect to the acceptance rate. Noting

$$|\log(\frac{\pi(\theta_{\mathcal{D}}^* | y) \pi_K(\theta | y)}{\pi(\theta_{\mathcal{D}} | y) \pi_K(\theta^* | y)})| \leq |\log(\pi(\theta_{\mathcal{D}}^* | y)) - \log(\pi_K(\theta^* | y))| + |\log(\pi(\theta_{\mathcal{D}} | y)) - \log(\pi_K(\theta | y))|, \quad (8)$$

it is sensible choose $\theta_{\mathcal{D}} \in \mathcal{D}$ to minimize the difference $Q(\theta_{\mathcal{D}}) = |\log(\pi(\theta_{\mathcal{D}} | y)) - \log(\pi_K(\theta | y))|$ towards 0 for each sample in the extrinsic posterior. Obviously, when the approximate $\theta \in \mathcal{D}$ exactly, the optimal projection would be the identity function; when $\theta \notin \mathcal{D}$, standard constrained optimization technique can be used.

Continuing the unit circle example, we first obtained posterior sample from $\pi_K(\theta | y)$ with $\lambda = 10^3$. The almost immediate drop to 0 in ACF indicates a rapid convergence to the target posterior $\pi_K(\theta | y)$. We then obtain $\hat{\theta}_{\mathcal{D}} = \underset{\theta_{\mathcal{D}}: \theta_{\mathcal{D}} = 1}{\operatorname{argmin}} |F' \theta_{\mathcal{D}} - \{F' \theta - \lambda(\theta' \theta - 1)^2\}|$ and construct the exact Markov chain. The acceptance rate is 97.9%.

3 Illustration

Example 1: Ordered Simplex

Example 2: Monotone Spline

Example 3: Orthonormal Gaussian Processes

4 Theory

Posterior propriety of $\pi_{0,\mathcal{R}}(\theta)\mathcal{E}(\theta)$, when $\pi_{0,\mathcal{D}}(\theta)$ is proper but $\pi_{0,\mathcal{R}}(\theta)$ is improper.

5 Application

6 Discussion

References

- Boyd, S. and L. Vandenberghe (2004). *Convex optimization*. Cambridge university press.
- Byrne, S. and M. Girolami (2013). Geodesic monte carlo on embedded manifolds. *Scandinavian Journal of Statistics* 40(4), 825–845.
- Gelfand, A. E., A. F. Smith, and T.-M. Lee (1992). Bayesian analysis of constrained parameter and truncated data problems using gibbs sampling. *Journal of the American Statistical Association* 87(418), 523–532.
- Girolami, M. and B. Calderhead (2011). Riemann manifold langevin and hamiltonian monte carlo methods. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 73(2), 123–214.
- Hoff, P. D. (2009). Simulation of the matrix bingham–von mises–fisher distribution, with applications to multivariate and relational data. *Journal of Computational and Graphical Statistics* 18(2), 438–456.
- Hoff, P. D. et al. (2016). Equivariant and scale-free tucker decomposition models. *Bayesian Analysis* 11(3), 627–648.
- Hoffman, M. D. and A. Gelman (2014). The no-u-turn sampler: adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research* 15(1), 1593–1623.
- Kelly, C. and J. Rice (1990). Monotone smoothing with application to dose-response curves and the assessment of synergism. *Biometrics*, 1071–1085.
- Khatri, C. and K. Mardia (1977). The von mises-fisher matrix distribution in orientation statistics. *Journal of the Royal Statistical Society. Series B (Methodological)*, 95–106.
- Lin, L. and D. B. Dunson (2014). Bayesian monotone regression using gaussian process projection. *Biometrika*.
- Lin, L., V. Rao, and D. B. Dunson (2016). Bayesian nonparametric inference on the stiefel manifold. *Statistica Sinica*.

- Lin, L., B. St Thomas, H. Zhu, and D. B. Dunson (2016). Extrinsic local regression on manifold-valued data. *Journal of the American Statistical Association* (just-accepted).
- Livingstone, S., M. Betancourt, S. Byrne, and M. Girolami (2016). On the geometric ergodicity of hamiltonian monte carlo. *arXiv preprint arXiv:1601.08057*.
- Murray, I., Z. Ghahramani, and D. MacKay (2012). Mcmc for doubly-intractable distributions. *arXiv preprint arXiv:1206.6848*.
- Neal, R. M. (2011). Mcmc using hamiltonian dynamics. *Handbook of Markov Chain Monte Carlo* 2, 113–162.
- Rao, V., L. Lin, and D. B. Dunson (2016). Data augmentation for models based on rejection sampling. *Biometrika* 103(2), 319–335.
- Uschmajew, A. (2010). Well-posedness of convex maximization problems on stiefel manifolds and orthogonal tensor product approximations. *Numerische Mathematik* 115(2), 309–331.
- Zhang, Y., Z. Ghahramani, A. J. Storkey, and C. A. Sutton (2012). Continuous relaxations for discrete hamiltonian monte carlo. In *Advances in Neural Information Processing Systems*, pp. 3194–3202.